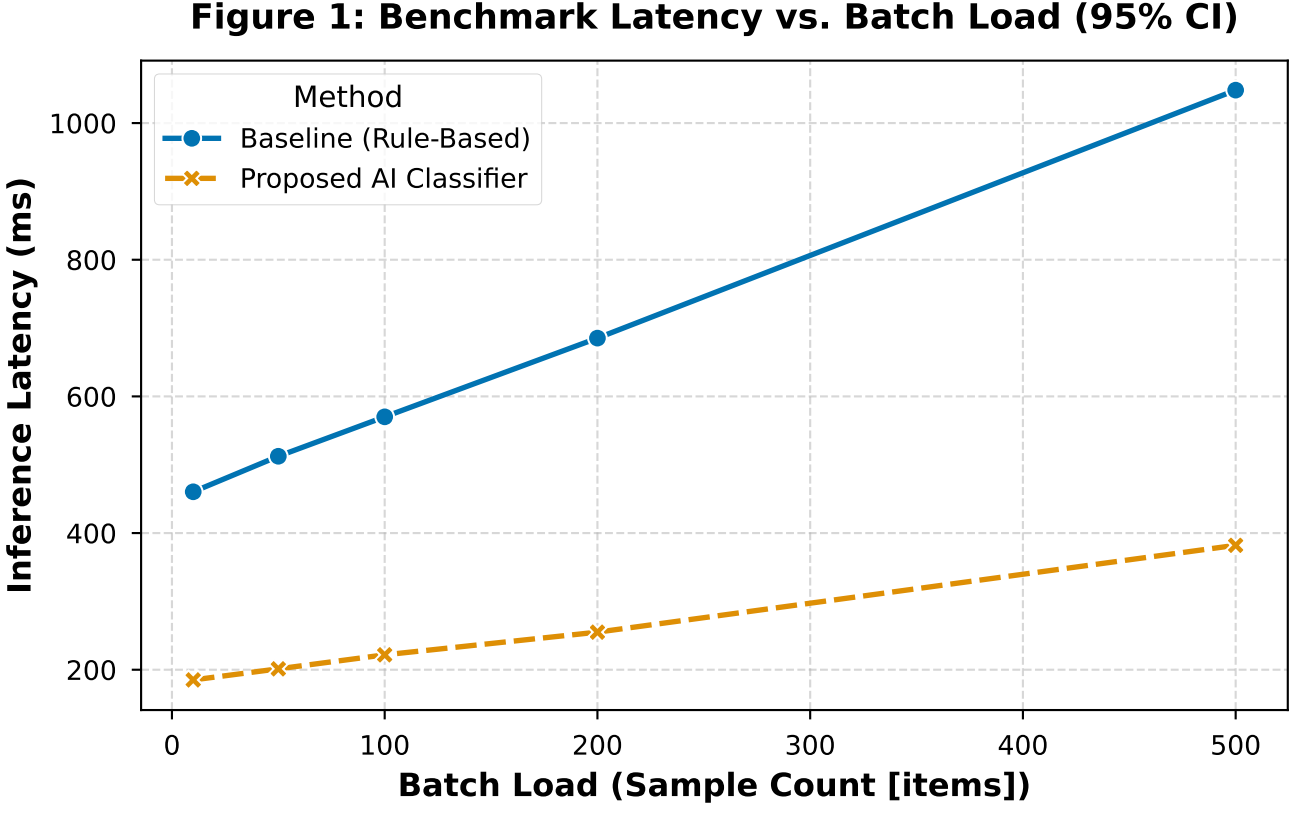
Figure 1: Benchmark Latency vs. Batch Load (95% CI)
Method
Baseline (Rule-Based)
Proposed AI Classifier
Inference Latency (ms)
1000
800
600
400
200
0
100
200
300
400
500
Batch Load (Sample Count [items])